

# CSC 112: Computer Operating Systems

## Lecture 3

### Synchronization

Department of Computer Science,  
Hofstra University

# Concurrency I

---

```
int x = 0;
```

```
//Thread T0  
for (int i=0; i<5; i++) {  
    x = x + 1;  
}
```

```
//Thread T1  
for (int j=0; j<5; j++) {  
    x = x + 2;  
}
```

- Consider two concurrent threads T0, T1, which access a shared variable x that has been initialized to 0. There is no mutex protection.
- Q: What are the minimum, maximum, and all possible values of x after the two threads have completed execution?

# Concurrency I Answer

---

```
int x = 0;
```

```
//Thread T0  
for (int i=0; i<5; i++) {  
    x = x + 1;  
}
```

```
//Thread T1  
for (int j=0; j<5; j++) {  
    x = x + 2;  
}
```

- ANS: Possible values of x after the two threads have completed execution: 5,...15. Min: 5. Max: 15.
  - The  $x=x+2$  statements can be “erased” by “sneaking in between” the load and store of an  $x=x+1$  statement, and vice versa. Each  $x=x+1$  statement can either do nothing (if erased by Thread T1) or increase x by 1. Each  $x=x+2$  statement can either do nothing (if erased by Thread T0) or increase x by 2. Since there are 5 of each type, and since x starts at 0, x has min 5 and max  $(5*1)+(5*2)=15$ . Possible values are 5, 6, 7,...15, e.g., If three increments from Thread T0 and two increments from Thread T1 are applied, then  $x=(3\times 1)+(2\times 2)=7$ .

## Concurrency II

---

- Consider three concurrent threads T1, T2, T3, which access a shared variable D that has been initialized to 100. There is no mutex protection. What are the minimum and maximum possible values of D after the three threads have completed execution?
- ANS:

```
//Initialization
int D=100;
//Thread T1
void main(){
D=D+20;
}
//Thread T2
void main(){
D=D-50;
}
//Thread T3
void main(){
D=D+10;
}
```

## Concurrency II Answer

---

- Min  $100 - 50 = 50$ , max  $100 + 20 + 10 = 130$
- Since each thread may read the value of int D, then write them in arbitrary order, overwriting each other's updates. Other possible results include  $100 + 10 = 110$ ,  $100 + 20 = 120$ ,  $100 + 20 - 50 = 70$ , and so on.

## Locks: Loads/Stores

- This implementation does not ensure mutual exclusion, since both threads may grab the lock:
- After Thread 1 reads `flag==0` and exits the `while(flag)` spin-wait loop, it is preempted/interrupted by Thread 2, which also reads `flag==0` and exits the `while(flag)` spin-wait loop. Then both threads set `flag=1` and enter the critical section.
- Root cause: Lock is not an atomic operation!

```
1 typedef struct __lock_t { int flag; } lock_t;
2
3 void init(lock_t *mutex) {
4     // 0 -> lock is available, 1 -> held
5     mutex->flag = 0;
6 }
7
8 void lock(lock_t *mutex) {
9     while (mutex->flag == 1) // TEST the flag
10        ; // spin-wait (do nothing)
11     mutex->flag = 1; // now SET it!
12 }
13
14 void unlock(lock_t *mutex) {
15     mutex->flag = 0;
16 }
```

**flag = 0**

**Thread 1**

call lock()  
while (flag == 1);  
**interrupt: switch to Thread 2**

flag = 1; // set flag to 1 (too!)

**Thread 2**

call lock()  
while (flag == 1);  
flag = 1;  
**interrupt: switch to Thread 1**

# Mutual Exclusion I

---

```
Boolean S0, S1;  
S0=false, S1=false;
```

```
//Thread T0  
while (true) {  
    //Spin-waits if S0 == S1  
    while (S0 == S1);  
    //Critical section  
    S0 = S1;  
}
```

```
//Thread T1  
while (true) {  
    //Spin-waits if S0 != S1  
    while (S0 != S1);  
    //Critical section  
    S1 = !S0;  
}
```

- Does it achieve one of more of the correctness properties of a concurrent program:
  - Mutual exclusion: Only one thread in critical section at a time
  - Progress (deadlock-free): If several simultaneous requests, must allow one to proceed
  - Bounded waiting (starvation-free): Must eventually allow each waiting thread to enter
- Does it need the TestAndSet() instruction for atomic execution like the previous slide “Locks: Loads/Stores”?
- What is its major flaw?
- ANS:

# Mutual Exclusion I: Sample Execution

```
Boolean S0, S1;  
S0=false, S1=false;
```

```
//Thread T0  
while (true) {  
    while (S0 == S1);  
    //Critical section  
    S0 = S1;  
}
```

```
//Thread T1  
while (true) {  
    while (S0 != S1);  
    //Critical section  
    S1 = !S0;  
}
```

- T0 and T1 take turns to enter the critical section in strict alternation order .

|          | S0 | S1 |
|----------|----|----|
| Init     | F  | F  |
| T1 in CS |    |    |
|          | F  | T  |
| T0 in CS |    |    |
|          | T  | T  |
| T1 in CS |    |    |
|          | T  | F  |
| T0 in CS |    |    |
|          | F  | F  |

# Mutual Exclusion I Answer

---

- Mutual Exclusion: Achieved. Only one thread can enter its critical section at a time because the conditions  $S0 == S1$  and  $S0 != S1$  ensure that only one thread can proceed.
- Progress (Deadlock-Free): Achieved. It is not possible for each thread to be blocked forever waiting for each other.
- Bounded Waiting (Starvation-Free): Achieved. Both threads enter each one's critical section in strict alternation order, i.e., T0, T1, T0, T1...
- TestAndSet Instruction: Not required. The solution uses simple Boolean variables and logical operations. Each thread reads both S0 and S1, but T0 only updates S0 and T1 only updates S1. **They do not update the same variable, so no TestAndSet Instruction is needed.**
  - In the previous slide “Locks: Loads/Stores”, all threads read and update a single global shared flag variable, so CPU atomic instructions like TestAndSet is needed to ensure atomicity of (read+modify+write) of the shared flag variable.
- Major Flaw: The algorithm relies on both threads actively participating in strict alternation order, i.e., T0, T1, T0, T1... If one thread stops due to some program bug or crashing, or is delayed indefinitely, the other thread might be blocked forever, leading to a potential deadlock. This may not happen for the example of two simple while loops, but it is just for illustration, whereas in reality each thread may run a large program with complex control flow, and use these instructions as lock/unlock instructions.

# Mutual Exclusion II

---

```
Boolean flag[2];  
flag[0]=false, flag[1]=false;
```

```
//Thread T0  
while (true) {  
    flag[0] = true;  
    while (flag[1]==true);  
    /* Critical Section */  
    flag[0] = false;  
}
```

```
//Thread T1  
while (true) {  
    flag[1] = true;  
    while (flag[0]==true);  
    /* Critical Section */  
    flag[1] = false;  
}
```

- Does it achieve one or more of the correctness properties of a concurrent program:
  - Mutual exclusion: Only one thread in critical section at a time
  - Progress (deadlock-free): If several simultaneous requests, must allow one to proceed
  - Bounded waiting (starvation-free): Must eventually allow each waiting thread to enter
- ANS:

# Mutual Exclusion II: Sample Execution & Answer

```
Boolean flag[2];  
flag[0]=false, flag[1]=false;
```

```
//Thread T0  
while (true) {  
    flag[0] = true;  
    while (flag[1]==true);  
    /* Critical Section */  
    flag[0] = false;  
}
```

```
//Thread T1  
while (true) {  
    flag[1] = true;  
    while (flag[0]==true);  
    /* Critical Section */  
    flag[1] = false;  
}
```

- **Mutual Exclusion: Achieved.** The use of flags ensures that only one thread can enter its critical section at a time.
- **Progress (Deadlock-Free): Not satisfied.** If both threads set their flags simultaneously (T0 sets `flag[0] = true`; T1 sets `flag[1] = true`), then both will spin-wait in `while()` loop, blocking each other indefinitely, resulting in deadlock.
- **Bounded Waiting (Starvation-Free): Achieved.** One thread cannot repeatedly enter the CS and starve the other thread, if the other thread is waiting.

|          | Flag[0] | Flag[1] |
|----------|---------|---------|
| Init     | F       | F       |
| T0 tries | T       | F       |
| T0 in CS |         |         |
|          | F       | F       |
| T1 tries | F       | T       |
| T1 in CS |         |         |
|          | F       | F       |
| T0 tries | T       | F       |
| T1 tries | T       | T       |
| Deadlock |         |         |

# Mutual Exclusion III (Peterson's Solution)

---

```
Boolean flag[2];  
flag[0]=false, flag[1]=false;  
int turn = 0;
```

```
//Thread T0  
while (true) {  
    flag[0] = true;  
    turn = 1;  
    while (flag[1]==true && turn==1);  
    /* Critical Section */  
    flag[0] = false;  
}
```

```
//Thread T1  
while (true) {  
    flag[1] = true;  
    turn = 0;  
    while (flag[0]==true && turn==0);  
    /* Critical Section */  
    flag[1] = false;  
}
```

- Does it achieve one or more of the correctness properties of a concurrent program:
  - Mutual exclusion: Only one thread in critical section at a time
  - Progress (deadlock-free): If several simultaneous requests, must allow one to proceed
  - Bounded waiting (starvation-free): Must eventually allow each waiting thread to enter
- ANS:

# Mutual Exclusion III (Peterson's Solution): Sample Execution & Answer

- Each thread is **courteous**, and sets **turn = (index of the other thread)**, and spins if the other thread is ready to enter CS or in CS, and **turn == (index of the other thread)**.
- TestAndSet Instruction: Not required. T0 reads flag[1] and updates flag[0]; T1 reads flag[0] and updates flag[1]; **They do not update the same flag variable**, so no mutual exclusion is needed for flag variables. Both threads read and write the turn variable, but it is **write followed by read**, so no race condition, unlike previous slide with race due to read followed by write to flag variable.
- Mutual Exclusion: Achieved. The combination of flag and turn ensures that only one thread can enter its critical section at a time.
- Progress (Deadlock-Free): Achieved. The turn variable ensures that if both threads want to enter their critical sections, one will eventually proceed. It is not possible for each thread to be blocked forever waiting for each other.
- Bounded Waiting (Starvation-Free): Achieved. Each thread gets a fair chance to enter its critical section due to the alternation enforced by the turn variable.

|                 | Flag[0]              | Flag[1] | turn |
|-----------------|----------------------|---------|------|
| Init            | F                    | F       | 0    |
| T0 tries        | T                    | F       | 1    |
| T0 in CS        |                      |         |      |
|                 | F                    | F       | 1    |
| T1 tries        | F                    | T       | 0    |
| T1 in CS        |                      |         |      |
|                 | F                    | F       | 0    |
| T0 tries        | T                    | F       | 1    |
| <b>T1 tries</b> | T                    | T       | 0    |
| T0 in CS        | (T1 cannot enter CS) |         |      |
|                 | F                    | T       | 0    |
| T0 tries        | T                    | T       | 1    |
| <b>T1 in CS</b> | (T0 cannot enter CS) |         |      |
|                 | T                    | F       | 1    |

# Mutual Exclusion III (Peterson's Solution Variation)

```
Boolean flag[2];  
flag[0]=false, flag[1]=false;  
int turn = 0;
```

```
//Thread T0  
while (true) {  
    flag[0] = true;  
    turn = 0;  
    while (flag[1]==true && turn==1);  
    /* Critical Section */  
    flag[0] = false;  
}
```

```
//Thread T1  
while (true) {  
    flag[1] = true;  
    turn = 1;  
    while (flag[0]==true && turn==0);  
    /* Critical Section */  
    flag[1] = false;  
}
```

- Does it achieve one or more of the correctness properties of a concurrent program:
  - Mutual exclusion: Only one thread in critical section at a time
  - Progress (deadlock-free): If several simultaneous requests, must allow one to proceed
  - Bounded waiting (starvation-free): Must eventually allow each waiting thread to enter
- ANS:

# Mutual Exclusion III (Peterson's Solution Variation)

## Sample Execution & Answer

- This variation is similar to Peterson's Solution but with an incorrect implementation of the turn variable. Each thread is **selfish**, and sets **turn = (index of self thread)**, and spin-waits if the other thread is ready to enter CS or in CS, and **turn == (index of the other thread)**.
- Mutual Exclusion: Achieved. Only one thread can enter its critical section at a time due to the conditions on flag and turn.
- Progress (Deadlock-Free): Achieved. It is not possible for each thread to be blocked forever waiting for each other.
- Bounded Waiting (Starvation-Free): Not satisfied. A thread may be indefinitely delayed if the other repeatedly sets its flag and does not allow alternation via the turn variable, i.e., one thread can repeatedly enter the CS and starve the other thread, which will spin-wait at the while() loop forever.
- TestAndSet Instruction: Not required.
- Major Flaw: Incorrect handling of the turn variable leads to potential livelock or starvation.

|                           | Flag[0] | Flag[1] | turn |
|---------------------------|---------|---------|------|
| Init                      | F       | F       | 0    |
| T0 tries                  | T       | F       | 0    |
| T0 in CS                  |         |         |      |
|                           | F       | F       | 0    |
| T1 tries                  | F       | T       | 1    |
| T1 in CS                  |         |         |      |
|                           | F       | F       | 1    |
| T0 tries                  | T       | F       | 0    |
| T1 tries                  | T       | T       | 1    |
| T1 in CS                  |         |         |      |
|                           | T       | F       | 1    |
| T1 tries                  | T       | T       | 1    |
| T1 in CS                  |         |         |      |
|                           | T       | F       | 1    |
| T0 experiences starvation |         |         |      |

# Race Conditions

---

Consider the two threads each executing t1 and t2. Values of shared variables y and z are initialized to 0

```
int y=0, z=0;
```

```
1 t1() {  
2     int x;  
3     x = y + z;  
4 }
```

```
1 t2() {  
2     y = 1;  
3     z = 2;  
4 }
```

Q. Give all possible final values for x and the corresponding order of execution of instructions in t1 and t2.

- 1) t1 runs to the end first; then t2 runs to the end:  $x = 0+0 = 0$
- 2) t2 to line 2; then t1 to the end; then t2 to the end:  $x = 1+0 = 1$
- 3) t2 to the end; then t1 to the end:  $x = 1+2 = 3$

*Are there other possibilities giving additional values?*

# Race Conditions

---

- Addition operation  $x=y+z$  consist of multiple machine instructions in assembly language:
  - A. fetch operand  $y$  into register  $r1$
  - B. fetch operand  $z$  into register  $r2$
  - C. Set register  $r3 = r1 + r2$
  - D. store  $r3$  in memory location of  $x$
- If a task switch to  $t2$  occurs between  $t1$ 's assembly instructions A and B; and then  $t2$  runs to completion before switching back to  $t1$ :
  - $y$  is read as 0 ( $t2$  didn't set  $y$  yet)
  - $z$  is read as 2 ( $t2$  sets  $z$  before execution instruction B of add. in  $t1$ )
  - the sum is then  $x = 0 + 2 = 2$

```
int y=0, z=0;
```

```
1 t1() {  
2     int x;  
3     x = y + z;  
4 }
```

```
1 t2() {  
2     y = 1;  
3     z = 2;  
4 }
```

# Race Conditions

---

Q. Give a solution using semaphores.

Solution: we protect the addition  $x = y + z$  within a critical section, *using* a binary semaphore (mutex). This code guarantees that  $x$  can never have the value 1 or 2, possible values are  $x = 0, 3$

(We use a slightly different notation of `s.wait()/s.signal()` to denote `sem_wait(&s)` and `sem_post(&s)`. Note that a binary semaphore's `wait()` and `signal()` are equivalent to a mutex's `lock()` and `unlock()`, when used for mutual exclusion. But a semaphore (binary or not) may be used for general synchronization/sequencing, in addition to working as a mutex.)

```
int y=0, z=0;
semaphore s = 1;
```

```
1 t1() {
2     int x;
3     s.wait();
4     x = y + z;
5     s.signal();
6 }
```

```
1 t2() {
2     s.wait();
3     y = 1;
4     z = 2;
5     s.signal();
6 }
```

# Semaphores I

t1:

```
1 int t1() {
2     printf("w");
3     printf("d");
4 }
```

t2:

```
1 int t2() {
2     printf("o");
3     printf("r");
4     printf("l");
5     printf("e");
6 }
```

Q. Use semaphores and insert wait/signal calls into the two threads so that only "wordle" is printed.

semaphore s1=1, s2=0

```
1 int t1() {
2     s1.wait();
3     printf("w");
4     s2.signal();
5     s1.wait();
6     printf("d");
7     s2.signal();
8 }
```

```
1 int t2() {
2     s2.wait();
3     printf("o");
4     printf("r");
5     s1.signal();
6     s2.wait();
7     printf("l");
8     printf("e");
9 }
```

- t1 has to run first to print "w", so s1 should be initialized to 1.
- t2 has to wait until the "w" has been printed by t1, then it is woken up by t1 calling s2.signal(), so s2 should be initialized to 0.
- Note that we reuse s1 and s2 in the two threads, but we can also declare additional semaphores s3 and s4 for use in the last two wait-signal pairs

## Semaphores II

---

- The following three functions of a program f1(), f2(), f3() run in separate threads each and print some prime numbers. All three threads are ready to run at the same time. Use synchronization using the semaphores S1, S2 and S3 and wait/signal operations on the semaphores to ensure that the program outputs the prime numbers in increasing order (2, 3, 5, 7, 11, 13).

```
Semaphore S1=0;
Semaphore S2=0;
Semaphore S3=0;
f1() {
    printf("3");
    printf("5");
}
f2() {
    printf("2");
    printf("13");
}
f3() {
    printf("7");
    printf("11");
}
```

## Semaphores II Solution

- Solution 1 (left): With initial values of all semaphores = 0, only f2 can run, prints 2, signals S1 and then waits for S2. S1.signal() starts f1, which was waiting for S1 and can now print 3 and 5 and then signal S3. S3.signal() now starts f3, which prints 7 and 11 and signals S2. This returns execution to f2, which can then finally print 13.
- Solution 2 (right): S2 has initial value 1, so f2 calls S2.wait() and runs first. The rest of the same as Solution 1. You can see that initializing S2=0 has the same effect as initializing S2=1 and let f2 call S2.wait() first. So Solution 1 is better with one less call to wait().

```
semaphore S1=0;
semaphore S2=0;
semaphore S3=0;
f1 () {
    S1.wait();
    printf("3");
    printf("5");
    S3.signal();
}
f2 () {
    printf("2");
    S1.signal();
    S2.wait();
    printf("13");
}
f3 () {
    S3.wait();
    printf("7");
    printf("11");
    S2.signal();
}
```

```
semaphore S1=0;
semaphore S2=1;
semaphore S3=0;
f1 () {
    S1.wait();
    printf("3");
    printf("5");
    S3.signal();
}
f2 () {
    S2.wait();
    printf("2");
    S1.signal();
    S2.wait();
    printf("13");
}
f3 () {
    S3.wait();
    printf("7");
    printf("11");
    S2.signal();
}
```

# Semaphores III

```
semaphore s_a=0, s_b=0, s_c=0;
```

```
1 int t1() {  
2     while(1) {  
3         printf("A");  
4         s_c.signal();  
5         s_a.wait();  
6     }  
7 }
```

```
1 int t2() {  
2     while(1) {  
3         printf("B");  
4         s_c.signal();  
5         s_b.wait();  
6     }  
7 }
```

```
1 int t3() {  
2     while(1) {  
3         s_c.wait();  
4         s_c.wait();  
5         printf("C");  
6         s_a.signal();  
7         s_b.signal();  
8     }  
9 }
```

Q. Which strings can be output when running the 3 threads in parallel?

- Either t1 or t2 could start first, so the first letter can be A or B
- Then both t1 and t2 signal s\_c, only after both have signalled s\_c, t3 can start and print C
- t3 signals s\_a and s\_b, which start in arbitrary order again
  - Even though t1 is woken up from its s\_a.wait() before t2 is woken up from its s\_b.wait(), both threads go into the ready queue, and the OS scheduler may schedule either thread to run before the other, depending on the scheduling algorithm and other code in each thread and other threads
- Accordingly, the output is a regular expression ((AB|BA)C)+
  - Print A or B in arbitrary order, then print C, then the process repeats

# Deadlocks I

```
//Initialization
int x=0, y=0, z=0;
semaphore lock1=1, lock2=1;
```

|                                                                                                                                                      |                                                                                                                                                      |
|------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------|
| <pre>1 int t1() { 2   z = z + 2; 3   lock1.wait(); 4   x = x + 2; 5   lock2.wait(); 6   lock1.signal(); 7   y = y + 2; 8   lock2.signal(); 9 }</pre> | <pre>1 int t2() { 2   lock2.wait(); 3   y = y + 1; 4   lock1.wait(); 5   x = x + 1; 6   lock1.signal(); 7   lock2.signal(); 8   z = z + 1; 9 }</pre> |
|------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------|

## Deadlock scenario 1:

- t2 runs first until line 2 (so lock2=0, lock1=1); switch to t1
- t1 starts and runs until line 3 (so lock1=0, lock2=0); back to t2
- t2 waits for lock1 in line 4; switch to t1, which waits for lock2 in line 5

Thread t1 grabs lock1 and requests lock 2; t2 grabs lock2 requests lock 1  
Deadlock: a circular waiting condition, where each thread grabs one lock and requests the other.  
To prevent deadlocks, every thread should acquire locks in the same order, e.g. both acquire lock1 before lock2, or both acquire lock2 before lock1

```
//Initialization
int x=0, y=0, z=0;
semaphore lock1=1, lock2=1;
```

|                                                                                                                                                      |                                                                                                                                                      |
|------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------|
| <pre>1 int t1() { 2   z = z + 2; 3   lock1.wait(); 4   x = x + 2; 5   lock2.wait(); 6   lock1.signal(); 7   y = y + 2; 8   lock2.signal(); 9 }</pre> | <pre>1 int t2() { 2   lock2.wait(); 3   y = y + 1; 4   lock1.wait(); 5   x = x + 1; 6   lock1.signal(); 7   lock2.signal(); 8   z = z + 1; 9 }</pre> |
|------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------|

## Deadlock scenario 2:

- t1 runs first until line 4 (so lock1=0, lock2=1); switch to t2
- t2 starts and runs until line 3 (so lock1=0, lock2=0); back to t1
- t1 waits for lock2 in line 5; switch to t2, waits for lock1 in line 4

# Deadlocks II

- Q. What are the possible values of x, y and z in the deadlock state?
- t1 runs until Line 5 lock2.wait() and t2 runs until Line 4 lock1.wait(), so x = 2, y = 1, z = 2
- Q. What are the possible values of x, y and z if the program finishes successfully without a deadlock?
- t1 runs first to the end, then t2 (or vice versa): x=3, y=3, z=3
- In t1, lock1.signal() sets lock1=1, lock2.signal() sets lock2=1, this exiting the critical sections protected by lock1 and lock2.
- Since Line 2 of t1 “z=z+2”, and Line 8 of t2 “z=z+1” are not protected within a critical section, a thread switch may occur in the middle of each line, and one of the updates to z may be lost:
  - t1’s update z=z+2 is lost: t2 Line 8 reads z=0; before z is written back; switch to t1 Line 2, run t1 to the end; switch to t2 Line 8, write back z=0+1=1.
  - Or, t2’s update z=z+1 is lost: t1 Line 2 reads z=0; before z is written back; switch to t2 Line 2, run t2 to the end; switch to t1 Line 2, write back z=0+2=2.

```
int x=0, y=0, z=0;
semaphore lock1=1, lock2=1;
```

```
1 int t1() {
2     z = z + 2;
3     lock1.wait();
4     x = x + 2;
5     lock2.wait();
6     lock1.signal();
7     y = y + 2;
8     lock2.signal();
9 }
```

```
1 int t2() {
2     lock2.wait();
3     y = y + 1;
4     lock1.wait();
5     x = x + 1;
6     lock1.signal();
7     lock2.signal();
8     z = z + 1;
9 }
```